


Introduction to

Chicago Graph Database
Meet-Up

Alex DeMarzi

Databases




About Me

Built the Neography Gem (Ruby Wrapper to the Neo4j REST API)
Playing with Neo4j since 10/2009




- My Blog: <http://maxdemarzi.com>
- Find me on Twitter: @maxdemarzi
- Email me: maxdemarzi@gmail.com
- GitHub: <http://github.com/maxdemarzi>



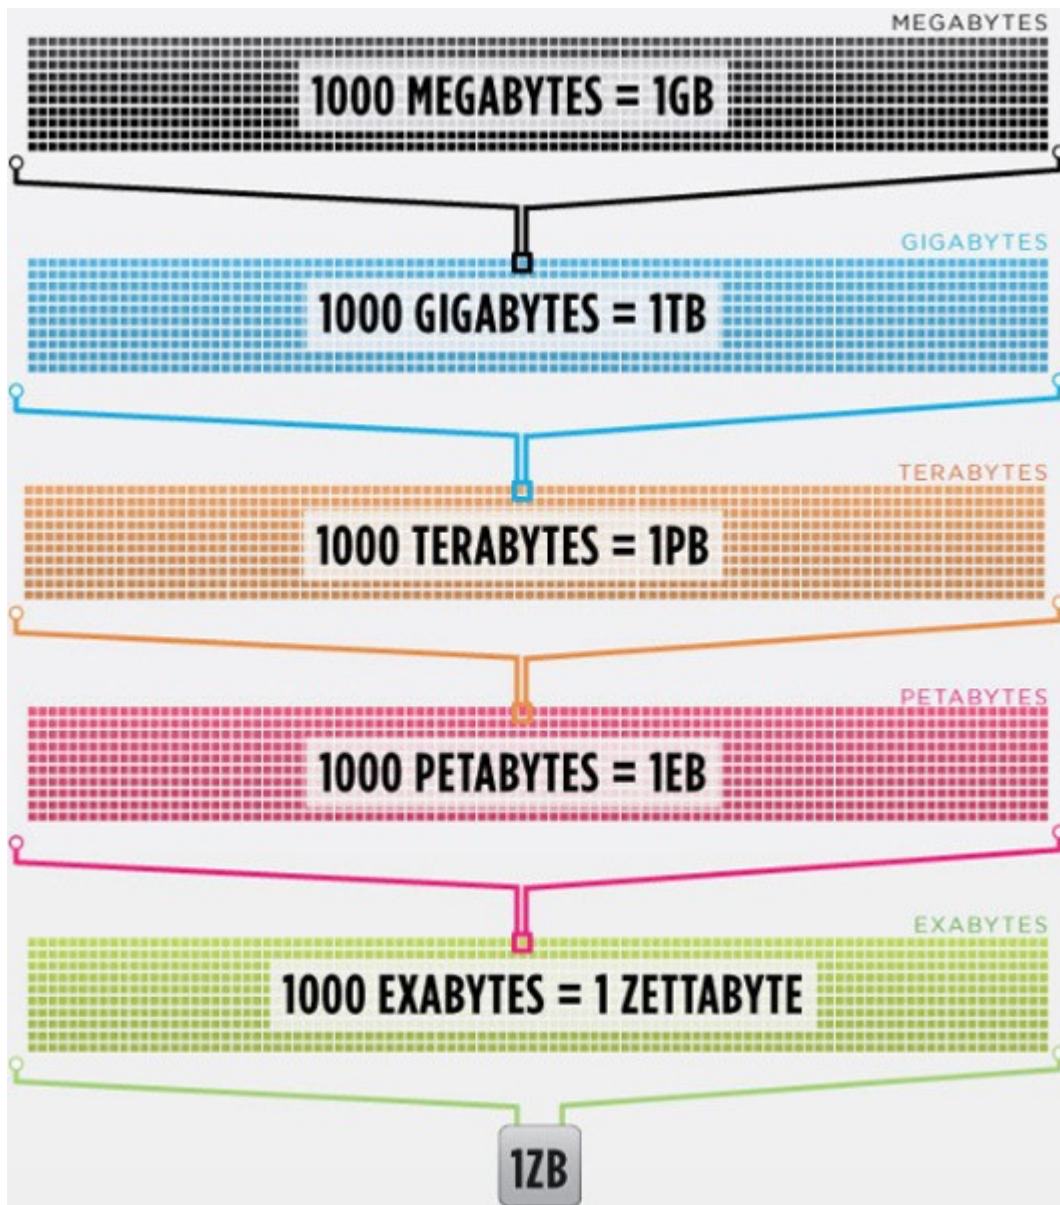
Agenda

- Trends in Data
- NOSQL
- What is a Graph?
- What is a Graph Database?
- What is Neo4j?



Trends in Data






Data is getting bigger:

“Every 2 days we create as much information as we did up to 2003”

– Eric Schmidt, Google


Data is more connected:

- Text (content)
- HyperText (added pointers)
- RSS (joined those pointers)
- Blogs (added pingbacks)
- Tagging (grouped related data)
- RDF (described connected data)
- GGG (content + pointers + relationships + descriptions)




Data is more Semi-Structured:

- If you tried to collect all the data of every movie ever made, how would you model it?
- Actors, Characters, Locations, Dates, Costs, Ratings, Showings, Ticket Sales, etc.



NOSQL

Not Only SQL




Less than 10% of the NOSQL Vendors




Key Value Stores

- Most Based on **Dynamo**: Amazon Highly Available Key-Value Store
- Data Model:
 - Global key-value mapping
 - Big scalable HashMap
 - Highly fault tolerant (typically)
- Examples:
 - Redis, Riak, Voldemort




Key Value Stores: Pros and Cons

- Pros:
 - Simple data model
 - Scalable
- Cons
 - Create your own “foreign keys”
 - Poor for complex data




Column Family

- Most Based on **BigTable**: Google's Distributed Storage System for Structured Data
- Data Model:
 - A big table, with column families
 - Map Reduce for querying/processing
- Examples:
 - HBase, HyperTable, Cassandra




Column Family: Pros and Cons

- Pros:
 - Supports Simi-Structured Data
 - Naturally Indexed (columns)
 - Scalable
- Cons
 - Poor for interconnected data




Document Databases

- Data Model:
 - A collection of documents
 - A document is a key value collection
 - Index-centric, lots of map-reduce
- Examples:
 - CouchDB, MongoDB




Document Databases: Pros and Cons

- Pros:
 - Simple, powerful data model
 - Scalable
- Cons
 - Poor for interconnected data
 - Query model limited to keys and indexes
 - Map reduce for larger queries




Graph Databases

- Data Model:
 - Nodes and Relationships
- Examples:
 - Neo4j, OrientDB, InfiniteGraph, AllegroGraph

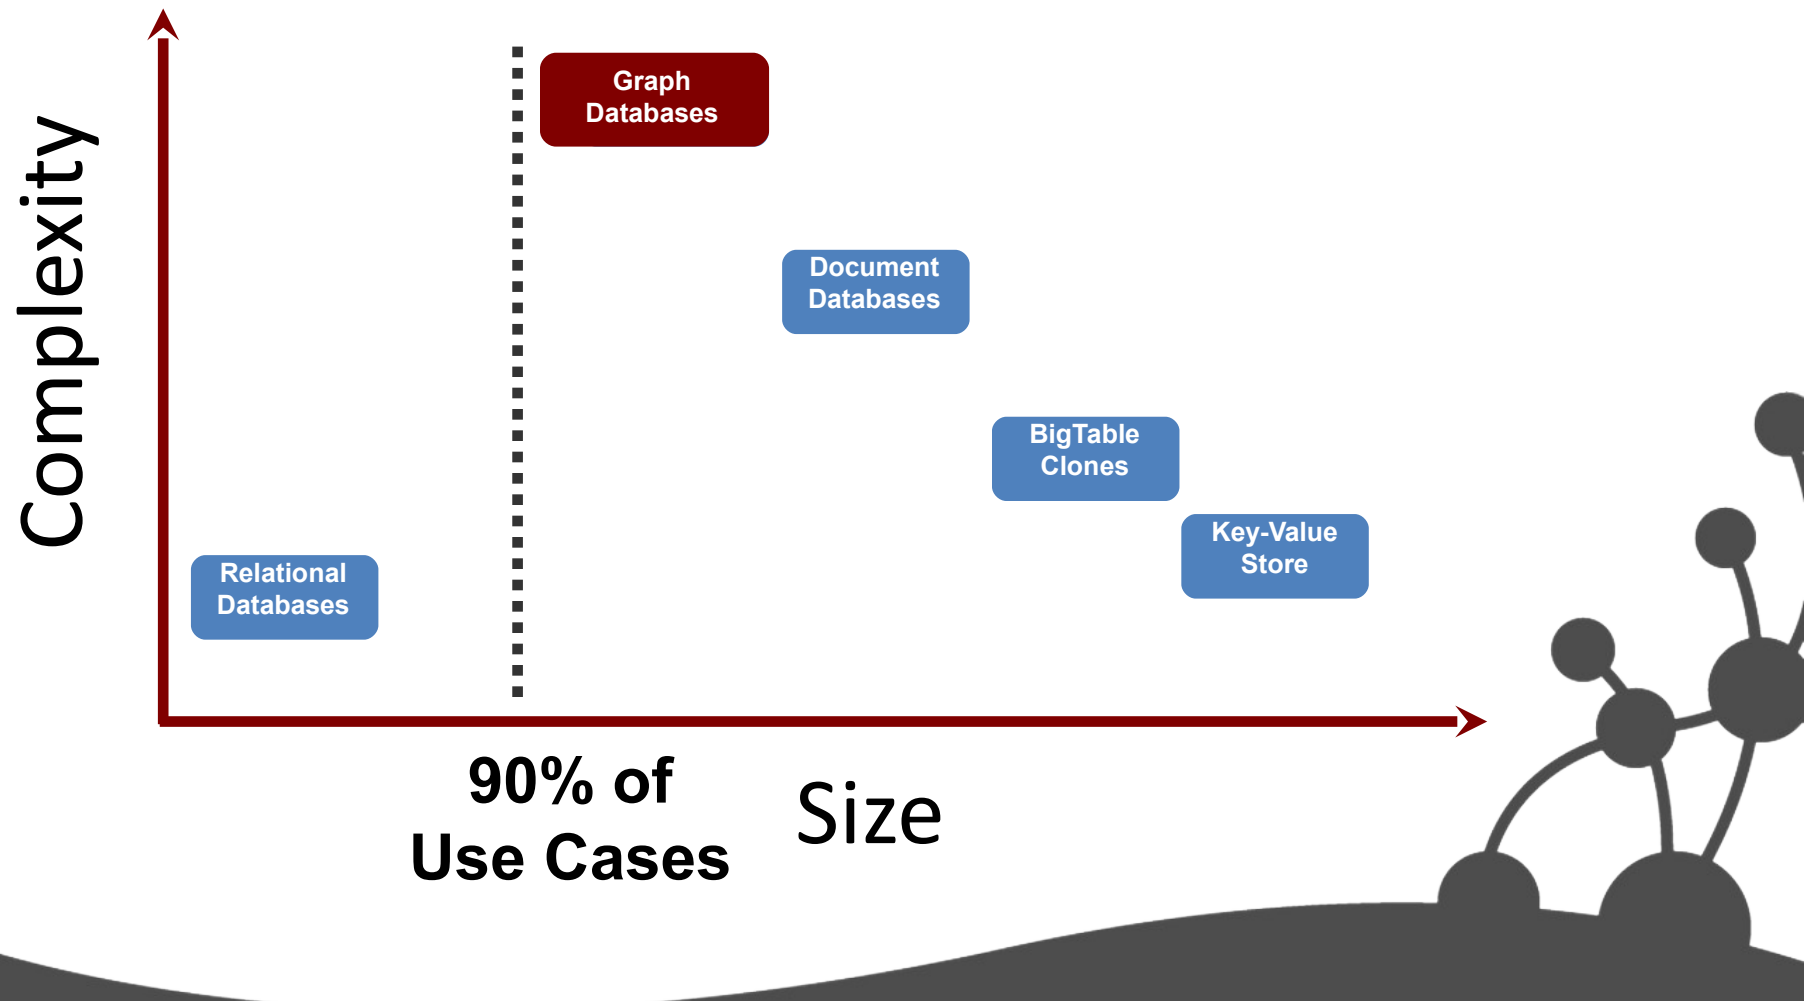


Graph Databases: Pros and Cons


- Pros:
 - Powerful data model, as general as RDBMS
 - Connected data locally indexed
 - Easy to query
- Cons
 - Sharding (lots of people working on this)
 - Scales UP reasonably well
 - Requires rewiring your brain



Living in a NOSQL World



What is a Graph?



What is a Graph?


- An abstract representation of a set of objects where some pairs are connected by links.



Object (Vertex, Node)




Link (Edge, Arc, Relationship)




Different Kinds of Graphs


- Undirected Graph




- Directed Graph




- Pseudo Graph



- Multi Graph




- Hyper Graph




More Kinds of Graphs


- Weighted Graph



- Labeled Graph




- Property Graph




What is a Graph Database?

- A database with an explicit graph structure
- Each node knows its adjacent nodes
- As the number of nodes increases, the cost of a local step (or hop) remains the same
- Plus an Index for lookups




Compared to Relational Databases

Optimized for aggregation



Optimized for connections




Compared to Key Value Stores

Optimized for simple look-ups



Optimized for traversing connected data




Compared to Key Value Stores


Optimized for “trees” of data



Optimized for seeing the forest and the trees, and the branches, and the trunks




What is Neo4j?




What is Neo4j?

- A Graph Database + Lucene Index
- Property Graph
- Full ACID (atomicity, consistency, isolation, durability)
- High Availability (with Enterprise Edition)
- 32 Billion Nodes, 32 Billion Relationships, 64 Billion Properties
- Embedded Server

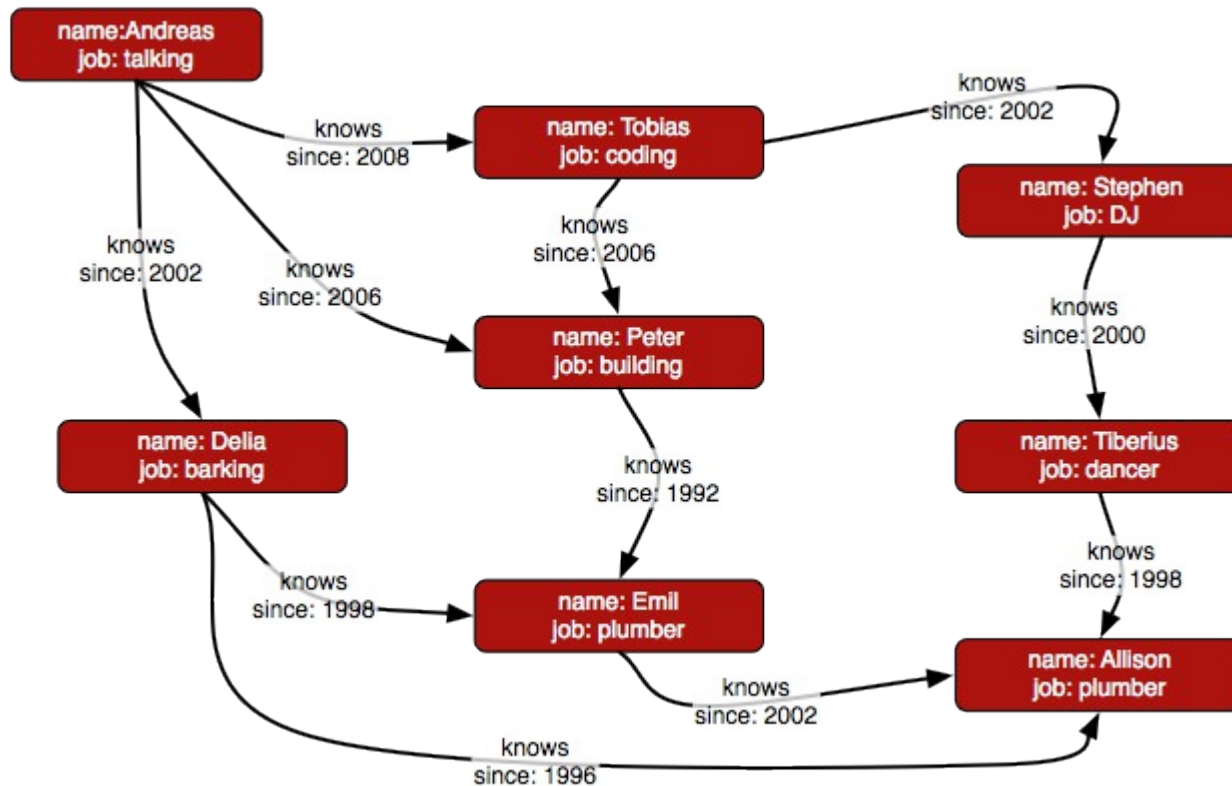


Good For

- Highly connected data (social networks)
- Recommendations (e-commerce)
- Path Finding (how do I know you?)
- A* (Least Cost path)
- Data First Schema (bottom-up, but you still need to design)

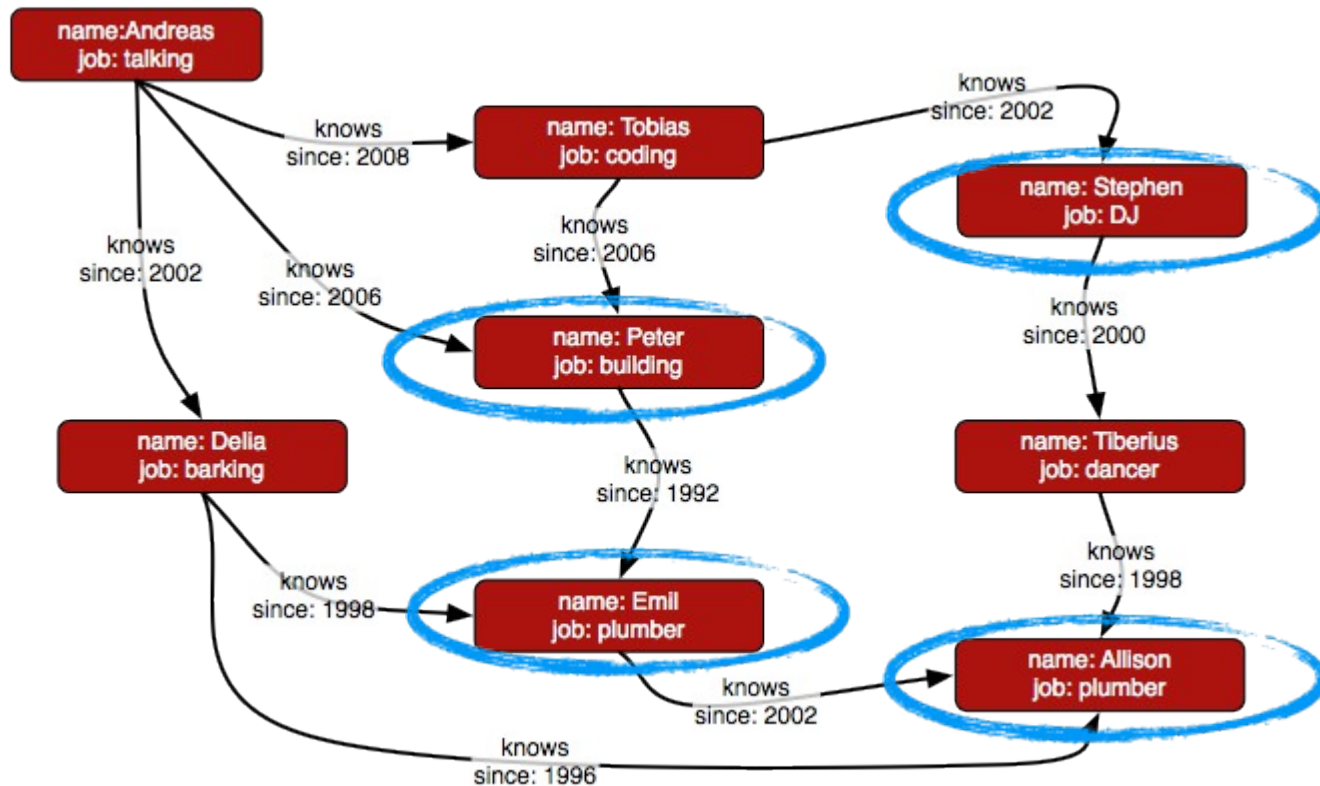


Property Graph



```
// then traverse to find results  
start n=(people-index, name, "Andreas")  
match (n)--()--(foaf) return foaf
```


n →



Cypher

Pattern Matching Query Language (like SQL for graphs)

```
// get node 0  
start a=(0) return a  
  
// traverse from node 1  
start a=(1) match (a)-->(b) return b  
  
// return friends of friends  
start a=(1) match (a)--()--(c) return c
```




Gremlin

A Graph Scripting DSL (groovy-based)

```
// get node 0
g.v(0)

// nodes with incoming relationship
g.v(0).in


// outgoing "KNOWS" relationship
g.v(0).out("KNOWS")
```




If you've ever

- Joined more than 7 tables together
- Modeled a graph in a table
- Written a recursive CTE
- Tried to write some crazy stored procedure with multiple recursive self and inner joins

You should use Neo4j





```
name: "Canada"
```

```
languages_spoken: "[ 'English', 'French' ]"
```

.....

spoken_in

```
language: "English"
```

spoken_in

```
name: "USA"
```


```
name: "Canada"
```

spoken_in


```
language: "French"
```

spoken_in

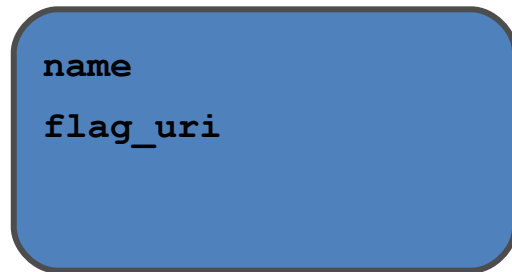
```
name: "France"
```



Country

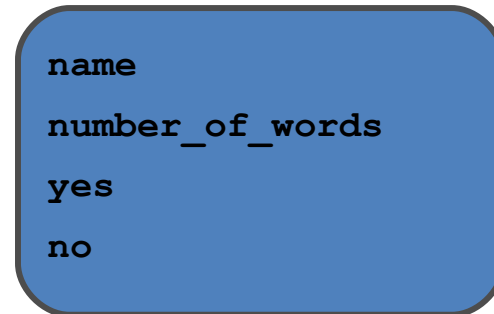


Country



SPEAKS

Language



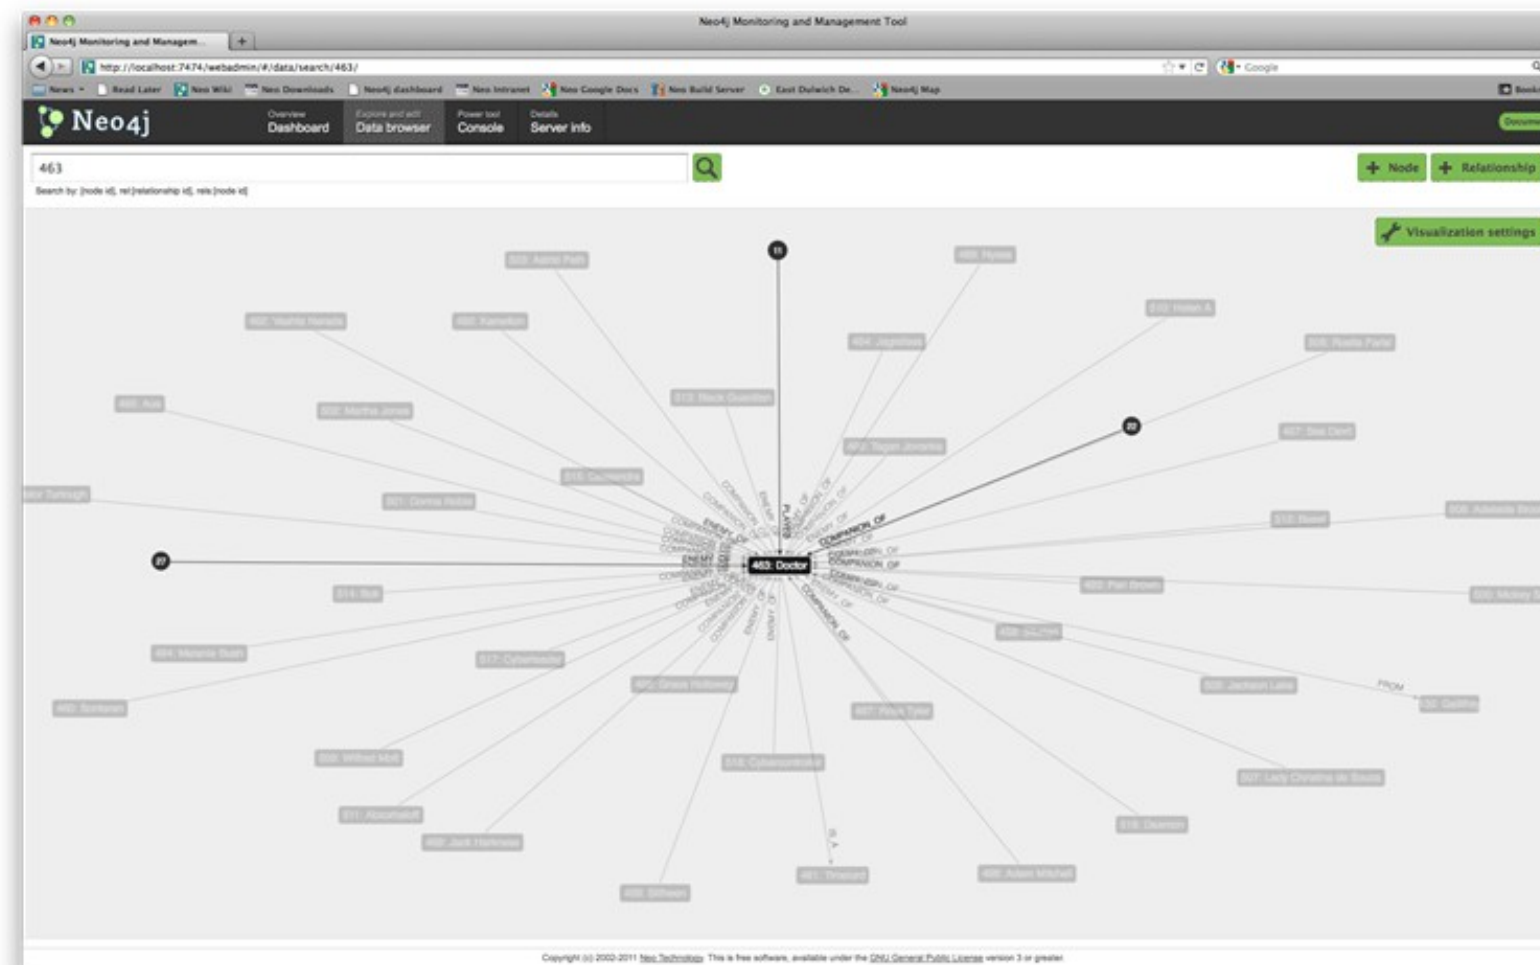
Currency



USES_CURRENCY



Neo4j Data Browser



Neo4j Console

Overview

Dashboard

Explore and edit

Data browser

Power tool

Console

Details

Server info

Indexing overview

Index manager

Documentation

Gremlin

Cypher

HTTP

```
==> The Five Doctors
==> The Chase
==> The Space Museum
==> The Crusade
==> The Web Planet
==> The Dalek Invasion of Earth
==> Planet of Giants
==> The Reign of Terror
==> The Sensorites
==> The Aztecs
==> The Keys of Marinus
==> Marco Polo
==> The Edge of Destruction
==> The Daleks
==> An Unearthly Child
==> +-----+
==> | 415 rows, 55 ms |
==> +-----+

cypher> start doctor=(Characters, name, 'Doctor') match (doctor)
<-[:COMPANION_OF]-(companion)-[:APPEARED_IN]->
(episode) return companion.name, count(episode) order by count(episode) desc limit 5
cypher>
==> +-----+
==> | companion.name | count(episode) |
==> +-----+
==> | Rose Tyler | 30 |
==> | Sarah Jane Smith | 22 |
==> | Jamie McCrimmon | 21 |
==> | Amy Pond | 21 |
==> | Tegan Jovanka | 20 |
==> +-----+
==> | 5 rows, 24 ms |
==> +-----+

cypher> |
```

console.neo4j.org

Try it right now:

start n=node(*) match n-[r:LOVES]->m return n, type(r), m

Notice the two nodes in red, they are your result set.

Graph Setup:

```
start root=node(0) create Neo = {name:'Neo'}, Morpheus = {name: 'Morpheus'}, Trinity = {name: 'Trinity'}, Cypher = {name: 'Cypher'}, Smith = {name: 'Agent Smith'}, Architect = {name:'The Architect'}, root-[:ROOT]->Neo, Neo-[:KNOWS]->Morpheus, Neo-[:LOVES]->Trinity, Morpheus-[:KNOWS]->Trinity, Morpheus-[:KNOWS]->Cypher, Cypher-[:KNOWS]->Smith, Smith-[:CODED_BY]->Architect
```







> start n=node(*) match n-[r?]->m return n,type(r),m

n	type(r)	m
Node[0] {}	"ROOT"	Node[1] {name->"Neo"}
Node[1] {name->"Neo"}	"KNOWS"	Node[2] {name->"Morpheus"}
Node[1] {name->"Neo"}	"LOVES"	Node[3] {name->"Trinity"}
Node[2] {name->"Morpheus"}	"KNOWS"	Node[3] {name->"Trinity"}
Node[2] {name->"Morpheus"}	"KNOWS"	Node[4] {name->"Cypher"}
Node[3] {name->"Trinity"}		
Node[4] {name->"Cypher"}	"KNOWS"	Node[5] {name->"Agent Smith"}
Node[5] {name->"Agent Smith"}	"CODED_BY"	Node[6] {name->"The Architect"}
Node[6] {name->"The Architect"}		

9 rows
0 ms

```
start n=node(*) match n-[r?]->m return n,type(r),m
```


What does a Graph look like?



Questions?

?



Thank you!

<http://maxdemarzi.com>

